

# Ruby API Survival Guide

# How to survive your first day in production


by Miguel Teixeira & Bruno Coelho



2 weeks before  
going live ...

# ALL good!

Web transactions time ▾









# Watch the red line!

Web transactions time ▾



# Panic!





" We can do this.  
We know what we are doing! "



# WEBrick

- single-process
- single-threaded
- bad for slow connections

# Thin

- event-driven
- single-process
- good for slow connections

# Unicorn

- worker pool (multi-process)
- single-threaded
- vulnerable to slow clients

# Puma

- single-process
- multi-threaded
- good for slow clients, slow I/O



" ... just add more threads,  
it's Puma we are using! "







# First chapter

Tune your app



# Tune your app

- choose your app server
- choose number of processes/workers
- choose number of threads
- size of connection pool
- `preload_app!`
- ...





# Second chapter

Optimal number of workers (processes)  
and threads

# # workers ?

## Container Size - Master Worker Process

---

1 Worker Process




# # threads ?

- GIL on MRI
- I/O
- request queuing
- Heroku



# Result





# Third chapter

## Profile your application



# ruby-prof

- easy to use
- fast
- configurable
- builtin reports



# Call stack

Call tree for application **primes**

Generated on 2015-04-24 11:52:24 +0200 with options {:application=>"primes"}

**Threshold:**

1.0

Apply

Expand All

Collapse All

Show Help

Thread: 70140045951280, Fiber: 70140054192180 (100.00% ~ 0.0029680728912353516)

- 100.00% (100.00%) [PrintersTest#setup](#) [1 calls, 1 total]

- 99.82% (99.82%) [Object#run\\_primes](#) [1 calls, 1 total]

- 92.62% (92.78%) [Object#find\\_primes](#) [1 calls, 1 total]

+ 92.55% (99.93%) [Array#select](#) [1 calls, 1 total]


+ 6.64% (6.65%) [Object#make\\_random\\_array](#) [1 calls, 1 total]

# benchmark-ips

- number of iterations per second
- configurable warm up phase
- easy to read report



# Result





Wrapping up



# Action points

- make your app tunnable

- LOAD TEST YOUR APP !

- choose # process & # threads

- AGAIN ... LOAD TEST YOUR APP !

- profile your app

- AGAIN & AGAIN ... LOAD TEST YOUR APP !

# Thank you!

- Availability team for always being available!
- DevOps for providing us numerous of burning machines 🔥
- and our own awesome team!



# Add to Pocket:

- Deploying Rails Applications with the Puma Web Server (<https://devcenter.heroku.com/articles/deploying-rails-applications-with-the-puma-web-server>)
- Scaling Ruby Apps to 1000 Requests per Minute - A Beginner's Guide (<https://www.speedshop.co/2015/07/29/scaling-ruby-apps-to-1000-rpm.html>)
- Your app server config is wrong @ RailsConf 2017 (<https://github.com/speedshop/appserverrailsconf2017>)
- Scaling Rails to 125,000 Requests per Minute on Heroku (<https://zeemee.engineering/scaling-rails-to-125-000-requests-per-minute-on-heroku-b4128a10a769>)

# Questions?

